

Assignment -5

SURESH KUMAR .R

2024-10-03

```
library(ISLR)
library(MASS)
library(class)
library(boot)
library(glmnet)

## Loading required package: Matrix

## Loaded glmnet 4.1-8

library(leaps)
library(pls)

##
## Attaching package: 'pls'

## The following object is masked from 'package:stats':
##
##      loadings

boston=Boston
```

11) We will now try to predict per capita crime rate in the Boston data set.

- a) Try out some of the regression methods explored in this chapter, such as best subset selection, the lasso, ridge regression, and PCR. Present and discuss results for the approaches that you consider.

Linear model:

```
fit_Q11_lm_c6=lm(crim~.,data = boston)
summary(fit_Q11_lm_c6)

##
## Call:
## lm(formula = crim ~ ., data = boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9.924 -2.120 -0.353  1.019 75.051
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  17.033228   7.234903   2.354 0.018949 *
```

```
## zn          0.044855    0.018734    2.394 0.017025 *
## indus      -0.063855    0.083407   -0.766 0.444294
## chas       -0.749134    1.180147   -0.635 0.525867
## nox       -10.313535    5.275536   -1.955 0.051152 .
## rm         0.430131    0.612830    0.702 0.483089
## age        0.001452    0.017925    0.081 0.935488
## dis       -0.987176    0.281817   -3.503 0.000502 ***
## rad        0.588209    0.088049    6.680 6.46e-11 ***
## tax       -0.003780    0.005156   -0.733 0.463793
## ptratio   -0.271081    0.186450   -1.454 0.146611
## black     -0.007538    0.003673   -2.052 0.040702 *
## lstat      0.126211    0.075725    1.667 0.096208 .
## medv      -0.198887    0.060516   -3.287 0.001087 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.439 on 492 degrees of freedom
## Multiple R-squared:  0.454, Adjusted R-squared:  0.4396
## F-statistic: 31.47 on 13 and 492 DF, p-value: < 2.2e-16
```


- In this model We can see that the variable zn,dis,rad,black,medv are having relationship with response.

Subset selection - best:

```
bs_Q11_c6=regsubsets(crim~.,data = boston,nvmax = 13)
bs_Q11_c6_summary<-summary(bs_Q11_c6)
bs_Q11_c6_summary$adjr2

## [1] 0.3900489 0.4184935 0.4251977 0.4289661 0.4336665 0.4373321 0.4398956
## [8] 0.4416149 0.4425053 0.4420078 0.4413928 0.4407131 0.4395838

par(mfrow=c(1,3))
plot(bs_Q11_c6_summary$cp,xlab = "Number of variable",ylab="C_p",type = "l")
points(which.min(bs_Q11_c6_summary$cp),bs_Q11_c6_summary$cp[which.min(bs_Q11_c6_summary$cp)],col="pink",cex=3,pch=20)
plot(bs_Q11_c6_summary$bic,xlab = "Number of variable",ylab="BIC",type = "l")
points(which.min(bs_Q11_c6_summary$bic),bs_Q11_c6_summary$bic[which.min(bs_Q11_c6_summary$bic)],col="pink",cex=3,pch=20)
plot(bs_Q11_c6_summary$adjr2,xlab = "Number of variable",ylab="adjr2",type = "l")
points(which.max(bs_Q11_c6_summary$adjr2),bs_Q11_c6_summary$adjr2[which.max(bs_Q11_c6_summary$adjr2)],col="pink",cex=3,pch=20)
```



```
coef(bs_Q11_c6,which.min(bs_Q11_c6_summary$cp))
```

```
## (Intercept)          zn          nox          dis          rad
## 19.683127801  0.043293393 -12.753707757 -0.918318253  0.532616533
##      ptratio      black      lstat      medv
## -0.310540942 -0.007922426  0.110173124 -0.174207166
```

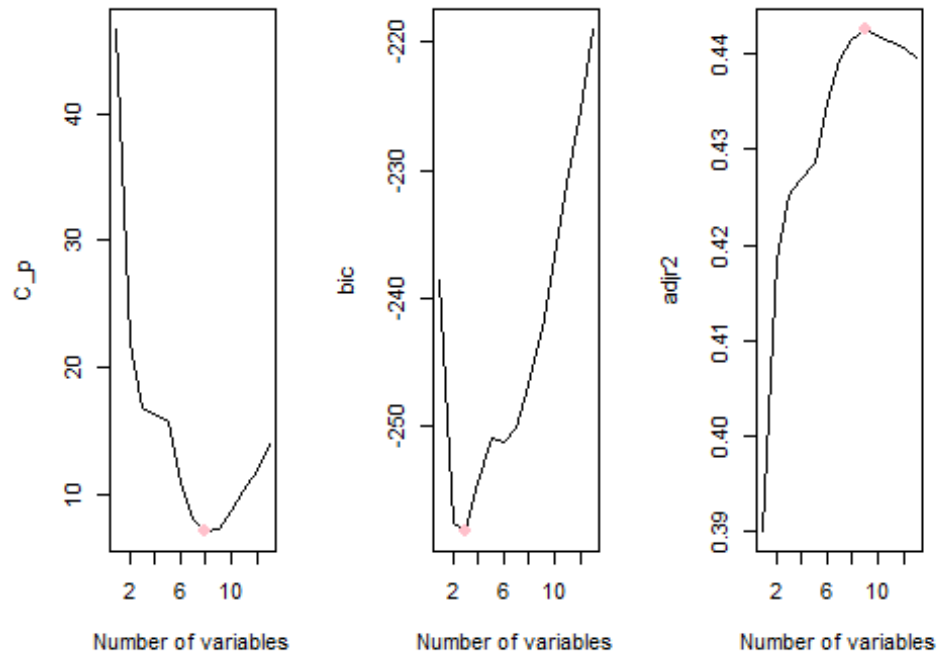
subset selection - forward:

```
fws_Q11_c6=regsubsets(crim~.,data = boston,nvmax = 13,method="forward")
fws_Q11_c6_summary<-summary(fws_Q11_c6)
fws_Q11_c6_summary$adjr2

## [1] 0.3900489 0.4184935 0.4251977 0.4268474 0.4286801 0.4350129 0.4395053
## [8] 0.4416149 0.4425053 0.4420078 0.4413928 0.4407131 0.4395838

par(mfrow=c(1,3))
plot(fws_Q11_c6_summary$cp,xlab = "Number of variables",ylab = "C_p",type =
"l")
points(which.min(fws_Q11_c6_summary$cp),fws_Q11_c6_summary$cp[which.min(fws_Q
11_c6_summary$cp)],col="pink",cex=2,pch=20)
plot(fws_Q11_c6_summary$bic,xlab = "Number of variables",ylab = "bic",type =
"l")
points(which.min(fws_Q11_c6_summary$bic),fws_Q11_c6_summary$bic[which.min(fws
_Q11_c6_summary$bic)],col="pink",cex=2,pch=20)
plot(fws_Q11_c6_summary$adjr2,xlab = "Number of variables",ylab =
"adjr2",type = "l")
```

```
points(which.max(fws_Q11_c6_summary$adjr2),fws_Q11_c6_summary$adjr2[which.max
(fws_Q11_c6_summary$adjr2)],col="pink",cex=2,pch=20)
```



```
coef(fws_Q11_c6,which.min(fws_Q11_c6_summary$cp))
```

```
## (Intercept)          zn          nox          dis          rad
## 19.683127801  0.043293393 -12.753707757 -0.918318253  0.532616533
##      ptratio         black         lstat         medv
## -0.310540942 -0.007922426  0.110173124 -0.174207166
```


subset selection - backward:

```
bws_Q11_c6=regsubsets(crim~.,data = boston,nvmax = 13,method="backward")
bws_Q11_c6_summary<-summary(bws_Q11_c6)
bws_Q11_c6_summary$outmat
```

```
##
## 1  ( 1 )  " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 2  ( 1 )  " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 3  ( 1 )  " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 4  ( 1 )  "*" " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 5  ( 1 )  "*" " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 6  ( 1 )  "*" " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 7  ( 1 )  "*" " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 8  ( 1 )  "*" " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 9  ( 1 )  "*" "*" " " " " " " " " " " " " " " " " " " " " " " " " "
## 10 ( 1 )  "*" "*" " " " " " " " " " " " " " " " " " " " " " " " " "
## 11 ( 1 )  "*" "*" " " " " " " " " " " " " " " " " " " " " " " " " "
```

```
## 12 ( 1 ) "*" "*" "*" "*" "*" " " " " "*" "*" "*" "*" "*" "*" "*"
## 13 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*"

par(mfrow=c(1,3))
plot(bws_Q11_c6_summary$cp,xlab = "Number of variables",ylab = "c_p",type =
"l")
points(which.min(bws_Q11_c6_summary$cp),bws_Q11_c6_summary$cp[which.min(bws_Q
11_c6_summary$cp)],col="blue",cex=3,pch=20)
plot(bws_Q11_c6_summary$bic,xlab = "Number of variables",ylab = "BIC",type =
"l")
points(which.min(bws_Q11_c6_summary$bic),bws_Q11_c6_summary$bic[which.min(bws
_Q11_c6_summary$bic)],col="blue",cex=3,pch=20)
plot(bws_Q11_c6_summary$adjr2,xlab = "Number of variables",ylab =
"adjr2",type = "l")
points(which.max(bws_Q11_c6_summary$adjr2),bws_Q11_c6_summary$adjr2[which.max
(bws_Q11_c6_summary$adjr2)],col="blue",cex=3,pch=20)
```



```
coef(bws_Q11_c6,which.max(bws_Q11_c6_summary$adjr2))

## (Intercept)          zn          indus          nox          dis
## 19.124636156  0.042788127 -0.099385948 -10.466490364 -1.002597606
##          rad          ptratio          black          lstat          medv
##  0.539503547 -0.270835584 -0.008003761  0.117805932 -0.180593877
```

- the subset selection the three method give similar variable's to use with response.

Regularization - ridge:

```

set.seed(2)
boston_matrix_crim<-model.matrix(crim~.,data = boston)[,-1]

ridge_c6_Q11=cv.glmnet(boston_matrix_crim,boston$crim,alpha=0)
bestlam_c6_ridge<-ridge_c6_Q11$lambda.min
bestlam_c6_ridge

## [1] 0.5374992

coef(ridge_c6_Q11,s=bestlam_c6_ridge)

## 14 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept)  9.063048666
## zn          0.033002416
## indus       -0.082046152
## chas        -0.737684583
## nox         -5.393098508
## rm          0.335972073
## age         0.001962473
## dis        -0.702123643
## rad         0.422779055
## tax         0.003400607
## ptratio    -0.135911587
## black      -0.008483285
## lstat       0.142613436
## medv       -0.139604127

```

- The ridge method say's tax,black,age this are near to zero,so this variable's can exclude.

Regularization - lasso:

```

set.seed(1)
lasso_c6_Q11=cv.glmnet(boston_matrix_crim,boston$crim,alpha=1)
bestlam_c6_lasso<-lasso_c6_Q11$lambda.min
bestlam_c6_lasso

## [1] 0.05630926

coef(lasso_c6_Q11,s=bestlam_c6_lasso)

## 14 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) 12.319178096
## zn          0.035726832
## indus       -0.068876055
## chas        -0.577832639
## nox         -6.631559478
## rm          0.208676938
## age         .
## dis        -0.768388825

```

```
## rad      0.512333871
## tax      .
## ptratio  -0.179631375
## black    -0.007551172
## lstat     0.124630014
## medv     -0.154550130
```

- In the lasso method it say's age and tax is exact zero.

Dimension Reduction - pcr:


```
set.seed(2)
pcr_c6_Q11=pcr(crim~.,data=boston,scale=TRUE,validation="CV")
pcr_c6_Q11summary<- summary(pcr_c6_Q11)

## Data:      X dimension: 506 13
## Y dimension: 506 1
## Fit method: svdpc
## Number of components considered: 13
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              8.61    7.229    7.227    6.814    6.799    6.795    6.794
## adjCV           8.61    7.225    7.222    6.807    6.789    6.788    6.787
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          6.787    6.654    6.664    6.673    6.676    6.651    6.573
## adjCV        6.780    6.645    6.656    6.664    6.666    6.639    6.562
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8
comps
## X          47.70    60.36    69.67    76.45    82.99    88.00    91.14
93.45
## crim       30.69    30.87    39.27    39.61    39.61    39.86    40.14
42.47
##      9 comps 10 comps 11 comps 12 comps 13 comps
## X          95.40    97.04    98.46    99.52    100.0
## crim       42.55    42.78    43.04    44.13    45.4

pcr_c6_Q11summary

## NULL

validationplot(pcr_c6_Q11,val.type = "MSEP")
```



```
loadingspcr_c6<-pcr_c6_Q11$loadings[,1:8]
loadingspcr_c6
```

##	Comp 1	Comp 2	Comp 3	Comp 4	Comp 5
## zn	-0.260386431	-0.12232275	-0.38678511	-0.372729217	0.118916506
## indus	0.344970045	0.11705301	0.01606480	-0.006169698	-0.021737885
## chas	-0.002537395	0.40537659	0.20238642	-0.691013712	-0.529946773
## nox	0.337133326	0.24651528	0.02428149	-0.059281996	0.195467829
## rm	-0.212422422	0.45501084	-0.33953869	0.271293829	-0.009667228
## age	0.309700710	0.24541638	0.20333926	0.097307974	0.148147177
## dis	-0.309928466	-0.34687275	-0.16446671	-0.202636506	-0.103749280
## rad	0.303520488	0.05063099	-0.47074769	0.006961260	-0.228603393
## tax	0.327873512	0.02246586	-0.41371570	-0.020610143	-0.161224582
## ptratio	0.214014884	-0.31923649	-0.08428970	0.316836569	-0.617413009
## black	-0.197245373	0.01096616	0.43281439	0.264196294	-0.372847153
## lstat	0.320591379	-0.21252009	0.14991413	-0.230935416	0.179841594
## medv	-0.274450582	0.45649295	-0.12343134	0.178606624	-0.052157658
##	Comp 6	Comp 7	Comp 8		
## zn	-0.41568749	0.31376763	0.40779707		
## indus	-0.14617630	-0.28073637	0.68536670		
## chas	0.16755735	0.04586120	-0.02611489		
## nox	-0.19106738	-0.09130262	0.06443069		
## rm	0.13455802	0.43665640	0.07685682		
## age	-0.03372481	0.59530163	0.01637673		
## dis	-0.02283442	0.10080440	-0.03153353		
## rad	-0.19077101	-0.05510114	-0.45621119		
## tax	-0.27651588	-0.11403113	-0.10347135		


```
## ptratio 0.27380093 0.24947434 0.29757016
## black -0.72392947 0.07846025 -0.08540903
## lstat -0.08494270 0.38780273 -0.19278149
## medv -0.04877358 -0.14919632 -0.01134994
```

- In pcr model say's that taking 3 or 8 component are best

Dimension Reduction - pls:


```
set.seed(2)
plsr_c6_Q11=plsr(crim~.,data=boston,scale=TRUE,validation="CV")
plsr_c6_Q11summary<- summary(plsr_c6_Q11)

## Data:      X dimension: 506 13
## Y dimension: 506 1
## Fit method: kernelpls
## Number of components considered: 13
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept) 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
## CV              8.61   7.054   6.702   6.621   6.607   6.605   6.597
## adjCV           8.61   7.049   6.695   6.611   6.597   6.593   6.583
##      7 comps 8 comps 9 comps 10 comps 11 comps 12 comps 13 comps
## CV        6.586   6.578   6.572   6.573   6.573   6.573   6.573
## adjCV      6.573   6.566   6.560   6.561   6.562   6.562   6.562
##
## TRAINING: % variance explained
##      1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps 8
comps
## X          47.27   56.79   61.38   71.13   76.41   79.78   83.99
86.27
## crim       34.32   41.81   44.03   44.58   44.94   45.24   45.33
45.38
##      9 comps 10 comps 11 comps 12 comps 13 comps
## X          88.5    91.32   96.56   98.26   100.0
## crim        45.4    45.40   45.40   45.40   45.4

plsr_c6_Q11summary

## NULL

validationplot(plsr_c6_Q11,val.type = "MSEP")
```



```
loadingsplsr_c6_Q11=plsr_c6_Q11$loadings[,1:2]
loadingsplsr_c6_Q11
```

```
##           Comp 1      Comp 2
## zn      -0.24335054  0.53815945
## indus    0.34770483 -0.17781631
## chas     -0.01277319 -0.20332787
## nox       0.34017136 -0.15637797
## rm       -0.20731864  0.25770067
## age       0.30404305 -0.32036861
## dis      -0.30411964  0.32217571
## rad       0.33246649  0.42195276
## tax       0.35402888  0.34100696
## ptratio  0.22217831  0.01782896
## black    -0.22006673 -0.34948686
## lstat     0.32535380 -0.10839180
## medv     -0.28015730  0.06836246
```

- In pls model say's that taking 2 component are best
- b) Propose a model (or set of models) that seem to perform well on this data set, and justify your answer. Make sure that you are evaluating model performance using validation set error, crossvalidation, or some other reasonable alternative, as opposed to using training error.

```
tr_bos=sample(nrow(boston),nrow(boston)*0.70)
```

```
tr_Q11b_bos=boston[tr_bos,]
```

```
te_Q11b_bos=boston[-tr_bos,]
```

Regularization - ridge:

```
set.seed(3)
tr_bos_matrix=model.matrix(crim~.,data = tr_Q11b_bos)[,-1]
te_bos_matrix=model.matrix(crim~.,data = te_Q11b_bos)[,-1]

Q_11b_ridge=cv.glmnet(tr_bos_matrix,tr_Q11b_bos$crim,alpha=0)
bestlam_Q11b_ridge=Q_11b_ridge$lambda.min

pred_Q11b_ridge=predict(Q_11b_ridge,s=bestlam_Q11b_ridge,newx =
te_bos_matrix)
test_error_Q11br=mean((te_Q11b_bos$crim- pred_Q11b_ridge)^2)

rmse_Q11b_ridge= sqrt(test_error_Q11br)
(rmse_Q11b_ridge/mean(te_Q11b_bos$crim))*100

## [1] 185.3637
```

Regularization - lasso:

```
Q_11b_lasso=cv.glmnet(tr_bos_matrix,tr_Q11b_bos$crim,alpha=1)
bestlam_Q11b_lasso=Q_11b_lasso$lambda.min

pred_Q11b_lasso=predict(Q_11b_lasso,s=bestlam_Q11b_lasso,newx =
te_bos_matrix)
test_error_Q11bl=mean((te_Q11b_bos$crim - pred_Q11b_lasso)^2)

rmse_Q11b_lasso= sqrt(test_error_Q11bl)
(rmse_Q11b_lasso/mean(te_Q11b_bos$crim))*100

## [1] 185.1713
```

- c) Does your chosen model involve all of the features in the data set? Why or why not?
- The chosen model does not involve all of the features in the data set, because some not statistically significant to response.